

Introduction to AntConc

I

臺灣台語語料庫

Ivo Niethammer Graduate Institute of Linguistics, NTU May 22, 2026

(1) 伊**講**伊欲來。

(2) 伊**說**伊欲來。

Which sounds more natural in Taiwanese Hokkien? Why?

What is a Corpus?

Corpus — from Latin: *body*

- A **systematically collected** body of real language use
- Stored and explored **computationally**
- Can be **general** (reference corpora) or **specialised** (built for a specific research question)
- Examples: British National Corpus (100M words), Common Crawl (petabytes), 臺灣台語語料庫

“The language looks rather different
when you look at a lot of it at once.”

— John Sinclair, corpus linguist

Why Use a Corpus?

- See **patterns** invisible when reading one text at a time
- **Zoom in and out** between individual examples and the whole dataset
- Avoid **cherry-picking** the corpus shows everything, not just what you expect
- Replace intuition with **evidence**
- Test hypotheses **systematically and at scale**

How Big is a Corpus?

Type	Size	Example
General reference	Billions of words	Common Crawl, BNC
Research corpus	Millions of words	COCA (520M words)
Specialised	Thousands of words	Our corpus today

The size of the corpus should match your research question.

Entry Points into a Corpus

Tool	What it shows
Frequency lists	most to least frequent tokens
Keyword lists	what is distinctive vs. a reference corpus?
Concordance lines (KWIC)	every occurrence in context
Collocates	what words appear nearby, more than chance?

Concordance Lines: KWIC

Key Word In Context — the core tool of corpus analysis

<i>Left context</i>	Hit	<i>Right context</i>
---------------------	------------	----------------------

人	講	將心比心
伊按呢	講	，無歹意
莫共人	講	！
你	講	甲一大拖，伊根本都聽袂落去
伊	講	話真粗魯

- Read **downwards** — same word across many texts
- **Sort** by left or right context to reveal patterns instantly

Collocates

Collocates: words that co-occur with a search term more often than chance

- Measured statistically: Log-Likelihood (is it significant?) and MI (how strong?)
- Reveals what grammatical and lexical company a word keeps
- Classic example: English *cause* collocates almost exclusively with negative nouns

cause problems, cause damage, cause concern. Never cause happiness

You would not notice this from introspection alone

Two words can share a meaning yet behave very differently in context.

Collocate analysis asks: what does each word attract, and are they the same?

Today: Our Corpus Exercise

Corpus	臺灣台語語料庫 (Taiwan Taigi Corpus)
Tool	AntConc
Data	200 sentences each for 講 and 說
Question	Are they really synonyms?
Methods	KWIC concordance, word list, collocates, keyword analysis

Get the Files



`github.com/ivoniet-ntu/antconc-taigi-crashcourse`

<code>kong_seg.txt</code>	講 corpus (200 sentences)
<code>sue_seg.txt</code>	說 corpus (200 sentences)
<code>antconc_visual_guide.pdf</code>	step-by-step loading guide